


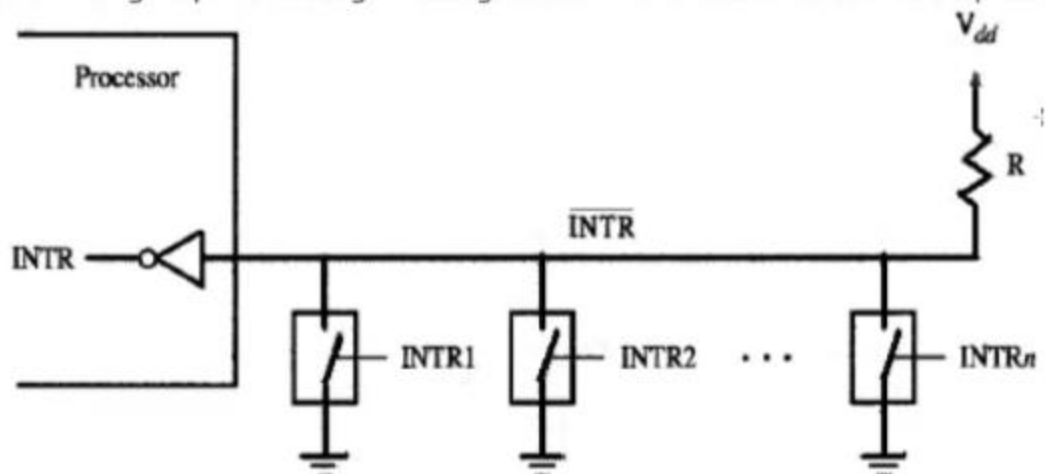
## Interrupts (contd..)



- Processor is executing the instruction located at address  $i$  when an interrupt occurs.
- Routine executed in response to an interrupt request is called the interrupt-service routine.
- When an interrupt occurs, control must be transferred to the interrupt service routine.
- But before transferring control, the current contents of the PC ( $i+1$ ), must be saved in a known location.
- This will enable the return-from-interrupt instruction to resume execution at  $i+1$ .
- Return address, or the contents of the PC are usually stored on the processor stack.

$$\text{INTR} = \text{INTR}_1 + \text{INTR}_2 + \dots + \text{INTR}_n$$

- A special gates known as open-collector or open-drain are used to drive the INTR line.
- The Output of the open collector control is equal to a switch to the ground that is
  - open when gates input is in "0" state and
  - closed when the gates input is in "1" state.
- Resistor R is called a **Pull-up Resistor** because it pulls the line voltage up to the high-voltage state when the switches are open.



**Figure 4.6** An equivalent circuit for an open-drain bus used to implement a common interrupt-request line.

### ENABLING & DISABLING INTERRUPTS

- All computers fundamentally should be able to enable and disable interruptions as desired.
- The problem of infinite loop occurs due to successive interruptions of active INTR signals.

## Interrupts (contd..)

- Interrupt-requests interrupt the execution of a program, and may alter the intended sequence of events:
  - Sometimes such alterations may be undesirable, and must not be allowed.
  - For example, the processor may not want to be interrupted by the same device while executing its interrupt-service routine.
- Processors generally provide the ability to enable and disable such interruptions as desired.
- One simple way is to provide machine instructions such as *Interrupt-enable* and *Interrupt-disable* for this purpose.
- To avoid interruption by the same device during the execution of an interrupt service routine:
  - First instruction of an interrupt service routine can be *Interrupt-disable*.
  - Last instruction of an interrupt service routine can be *Interrupt-enable*.

## Interrupts (contd..)

- Saving and restoring information can be done automatically by the processor or explicitly by program instructions.
- Saving and restoring registers involves memory transfers:
  - Increases the total execution time.
  - Increases the delay between the time an interrupt request is received, and the start of execution of the interrupt-service routine. This delay is called interrupt latency.
- In order to reduce the interrupt latency, most processors save only the minimal amount of information:
  - This minimal amount of information includes Program Counter and processor status registers.
- Any additional information that must be saved, must be saved explicitly by the program instructions at the beginning of the interrupt service routine.

## Interrupts (contd..)

- Interrupt-requests interrupt the execution of a program, and may alter the intended sequence of events:
  - Sometimes such alterations may be undesirable, and must not be allowed.
  - For example, the processor may not want to be interrupted by the same device while executing its interrupt-service routine.
- Processors generally provide the ability to enable and disable such interruptions as desired.
- One simple way is to provide machine instructions such as *Interrupt-enable* and *Interrupt-disable* for this purpose.
- To avoid interruption by the same device during the execution of an interrupt service routine:
  - First instruction of an interrupt service routine can be *Interrupt-disable*.
  - Last instruction of an interrupt service routine can be *Interrupt-enable*.

## Interrupts (contd..)

- Multiple I/O devices may be connected to the processor and the memory via a bus. Some or all of these devices may be capable of generating interrupt requests.
  - Each device operates independently, and hence no definite order can be imposed on how the devices generate interrupt requests?
- How does the processor know which device has generated an interrupt?
- How does the processor know which interrupt service routine needs to be executed?
- When the processor is executing an interrupt service routine for one device, can other device interrupt the processor?
- If two interrupt-requests are received simultaneously, then how to break the tie?





## **COMPUTER ORGANIZATION**

### **HANDLING MULTIPLE DEVICES**

- While handling multiple devices, the issues concerned are:
  - 1) How can the processor recognize the device requesting an interrupt?
  - 2) How can the processor obtain the starting address of the appropriate ISR?
  - 3) Should a device be allowed to interrupt the processor while another interrupt is being serviced?
  - 4) How should 2 or more simultaneous interrupt-requests be handled?

### **POLLING**

- Information needed to determine whether device is requesting interrupt is available in status-register
- Following condition-codes are used:
  - DIRQ → Interrupt-request for display.
  - KIRQ → Interrupt-request for keyboard.
  - KEN → keyboard enable.
  - DEN → Display Enable.
  - SIN, SOUT → status flags.
- For an input device, SIN status flag is used
  - SIN = 1 → when a character is entered by user.
  - SIN = 0 → when the character is read by processor.

Manasa R

meet.google.com is sharing your screen.

Stop sharing

hide

- 1) How can the processor recognize the device requesting an interrupt?
- 2) How can the processor obtain the starting address of the appropriate ISR?
- 3) Should a device be allowed to interrupt the processor while another interrupt is being serviced?
- 4) How should 2 or more simultaneous interrupt-requests be handled?

## POLLING

- Information needed to determine whether device is requesting interrupt is available in status-register
- Following condition-codes are used:
  - DIRQ → Interrupt-request for display.
  - KIRQ → Interrupt-request for keyboard.
  - KEN → keyboard enable.
  - DEN → Display Enable.
  - SIN, SOUT → status flags.
- For an input device, SIN status flag is used.
  - SIN = 1 → when a character is entered at the keyboard.
  - SIN = 0 → when the character is read by processor.
  - IRQ=1 → when a device raises an interrupt-requests (Figure 4.3).
- Simplest way to identify interrupting-device is to have ISR poll all devices connected to bus.
- The first device encountered with its IRQ bit set is serviced.
- After servicing first device, next requests may be serviced.
- **Advantage:** Simple & easy to implement.
- **Disadvantage:** More time spent polling IRQ bits of all devices.

DATAIN

DATAOUT

meet.google.com is sharing your screen.

Stop sharing

Hide



## HANDLING MULTIPLE DEVICES

- While handling multiple devices, the issues concerned are:
  - 1) How can the processor recognize the device requesting an interrupt?
  - 2) How can the processor obtain the starting address of the appropriate ISR?
  - 3) Should a device be allowed to interrupt the processor while another interrupt is being serviced?
  - 4) How should 2 or more simultaneous interrupt-requests be handled?

## POLLING

- Information needed to determine whether device is requesting interrupt is available in status-register
- Following condition-codes are used:
  - DIRQ → Interrupt-request for display.
  - KIRQ → Interrupt-request for keyboard.
  - KEN → keyboard enable.
  - DEN → Display Enable.
  - SIN, SOUT → status flags.
- For an input device, SIN status flag is used.
  - SIN = 1 → when a character is entered at the keyboard.
  - SIN = 0 → when the character is read by processor.
  - IRQ=1 → when a device raises an interrupt-requests (Figure 4.3).
- Simplest way to identify interrupting-device is to have ISR poll all devices connected to bus.
- The first device encountered with its IRQ bit set is serviced.
- After servicing first device, next requests may be serviced.
- **Advantage:** Simple & easy to implement.
- **Disadvantage:** More time spent polling IRQ bits of all devices.

## VECTORED INTERRUPTS

- A device requesting an interrupt identifies itself by sending a special-code to processor over bus.
- Then, the processor starts executing the ISR.
- The special-code indicates starting-address of ISR.
- The special-code length ranges from 4 to 8 bits.
- The location pointed to by the interrupting-device is used to store the starting address to ISR.
- The starting address to ISR is called the **interrupt vector**.
- Processor
  - loads interrupt-vector into PC &
  - executes appropriate ISR.
- When processor is ready to receive interrupt-vector code, it activates INTA line.
- Then, I/O-device responds by sending its interrupt-vector code & turning off the INTR signal.
- The interrupt vector also includes a new value for the Processor Status Register.

## CONTROLLING DEVICE REQUESTS

- Following condition-codes are used:
  - KEN → Keyboard Interrupt Enable.
  - DEN → Display Interrupt Enable.
  - KIRQ/DIRQ → Keyboard/Display unit requesting an interrupt.
- There are 2 independent methods for controlling interrupt-requests. (IE → interrupt-enable).

### 1) At Device-end

IE bit in a control-register determines whether device is allowed to generate an interrupt-request.

### 2) At Processor-end, Interrupt-request is determined by

- IE bit in the PS register or
- Priority structure

## INTERRUPT NESTING

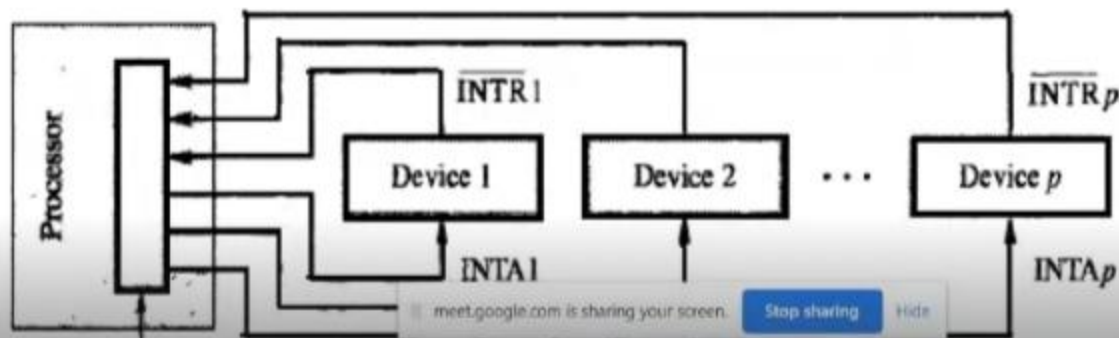
- A multiple-priority scheme is implemented by using separate INTR & INTA lines for each device
- Each INTR line is assigned a different priority-level (Figure 4.7).
- Priority-level of processor is the priority of program that is currently being executed.
- Processor accepts interrupts only from devices that have higher-priority than its own.
- At the time of execution of ISR for some device, priority of processor is raised to that of the device.
- Thus, interrupts from devices at the same level of priority or lower are disabled.

## Privileged Instruction

- Processor's priority is encoded in a few bits of PS word. (PS → Processor-Status).
- Encoded-bits can be changed by **Privileged Instructions** that write into PS.
- Privileged-instructions can be executed only while processor is running in **Supervisor Mode**.
- Processor is in supervisor-mode only when executing operating-system routines.

## Privileged Exception

- User program cannot
  - accidentally or intentionally change the priority of the processor &
  - disrupt the system-operation.
- An attempt to execute a privileged-instruction while in user-mode leads to a **Privileged Exception**.



## EXCEPTIONS

- An **interrupt** is an event that causes
  - execution of one program to be suspended &
  - execution of another program to begin.
- **Exception** refers to any event that causes an interruption. For ex: I/O interrupts.

### 1. Recovery from Errors

- These are techniques to ensure that all hardware components are operating properly.
- For ex: Many computers include an ECC in memory which allows detection of errors in stored-data.  
(ECC → Error Checking Code, ESR → Exception Service Routine).
- If an error occurs, control-hardware
  - detects the errors &
  - informs processor by raising an interrupt.
- When exception processing is initiated (as a result of errors), processor.
  - suspends program being executed &
  - starts an **ESR**. This routine takes appropriate action to recover from the error.

### 2. Debugging

- Debugger
  - is used to find errors in a program and
  - uses exceptions to provide 2 important facilities: i) Trace & ii) Breakpoints

#### i) Trace

- When a processor is operating in trace-mode, an exception occurs after execution of every instruction (using debugging-program as ESR).
- Debugging-program enables user to examine contents of registers, memory-locations and so on.
- On return from debugging-program,  
next instruction in program being debugged is executed